

# PageRank Algorithm

## Installation

```
1 1. Clone this repository
2 2. make pagerank
3 3. ./pagerank
```

## Running Perf, Benchmark & Validity

In order to run perf tests (outputted to `out`), timing and validity tests type:

```
1 1. ./test.sh
```

## Description

**NOTE: All versions of the pagerank methods can be found in `pagerank.c` just modify their names or the method being called in `main`.**

The data structure used for the implementation of the PageRank algorithm was a simple `struct`:

```
1 /**
2  * The struct to store the page and its score
3  */
4 struct page_score {
5     double score[2];
6     page* page;
7 };
```

Each would store an array of the scores rather than having a separate `previous` and `new_score` value for each struct. Alternatively (as discussed in [Score Array and Unnecessary Copying](#)), a score array was not the first approach but rather using two variables to store the old and new value with a copy after every iteration.

*Previous Struct*

```
1 struct page_score_2D {
2     double new_score;
3     double old_score;
4     page* page;
5 };
```

Furthermore each `page_score` struct would be initialised to hold the page of list in sequential order such that their indices correspond to the index in the `struct page`.

```
1 /**
2  * Initialise the values of the struct array of page scores
```

```

3      * @param plist, the list of pages
4      * @param npages, the number of pages
5      * @return the array of page score structs
6      */
7  struct page_score* init_pageranks(list* plist, int npages) {
8      struct page_score* page_scores = malloc(sizeof(struct page_score) * np
9      ages);
10     register double initial_value = 1/(double)(npages);
11     node* current = plist->head;
12     for (size_t i = 0; i < npages; i++) {
13         page_scores[i].page = current->page;
14         page_scores[i].score[0] = initial_value;
15         page_scores[i].score[1] = initial_value;
16         current = current->next;
17     }
18     return page_scores;
19 }

```

Hence, through this there was a way to hold the score at each iteration for each page without having to modify the provided node or page structs and without having to do an unnecessary search every time.

Afterwards the sequential implementation just required to loop until convergence. Within the loop every `page_score` is changed given its page `inlinks` and `outlinks`.

```

1  // Loop through until the convergence threshold is reached
2  while (diff > EPSILON) {
3      diff = 0.0;
4      for (size_t i = 0; i < npages; i++) {
5          page_scores[i].score[x] = dampening_value;
6          register double total = 0.0;
7
8          // Get the list of pages that inlink to this page
9          list* inlist = page_scores[i].page->inlinks;
10
11         // If null then add difference and continue looping
12         if (inlist == NULL || current->page->noutlinks == 0) {
13             diff += pow(page_scores[i].score[x] - page_scores[i].score[!x],
14             2);;
15             continue;
16         }
17         // Get the node to loop
18         node* current = inlist->head;
19
20         // Loop through the list
21         while (current != NULL) {
22             total += (page_scores[current->page->index].score[!x]) /
23             ((double)current->page->noutlinks);
24             current = current->next;
25         }
26         page_scores[i].score[x] += total * dampener; // Update the new score
27         diff += pow(page_scores[i].score[x] - page_scores[i].score[!x], 2);
28     }
29     x = (x + 1) % 2; // Update the value so we do not have to copy
30     diff = sqrt(diff); // Get the total difference
31 }

```

# Implementation and Benchmarks

## Score Array and Unnecessary Copying

Based on the previous part of the assignment it was decided that using some knowledge gained off a 3D array was useful in this scenario. Rather than having a copy loop which would copy all previous values of the iteration to the next.

```
1 for (size_t i = 0; i < npages; i++) {
2     page_scores[i].prev = page_scores[i].current;
3 }
```

Instead the `struct page_score` would contain a `double[2]` array and would alternate between 1 and 0 after each iteration.

```
1 for (size_t i = 0; i < npages; i++) {
2     page_scores[i].score[x] = dampening_value;
3
4     // CODE BELOW - removed for this example
5
6     while (current != NULL) {
7         total += (page_scores[current->page->index].score[!x]) /
8         ((double)current->page->noutlinks );
9         current = current->next;
10    }
11
12    page_scores[i].score[x] += total * dampener; // Update the new score
13    diff += (page_scores[i].score[x] - page_scores[i].score[!x]) *
14    (page_scores[i].score[x] - page_scores[i].score[!x]);
15 }
16
17 x = (x + 1) % 2; // Update the value so we do not have to copy
```

Specifically there would have had to have been something like this:

```
1 for (size_t i = 0; i < npages; i++) {
2     page_scores[i].old_score = page_scores[i].new_score;
3 }
```

These resulted in significantly longer runtimes:

```
1 ----- TIME TESTS -----
2 Test: sample.in
3
4 real    0m0.001s
5 user    0m0.001s
6 sys     0m0.000s
7 Test: test01.in
8
```

```
9 real 0m0.001s
10 user 0m0.000s
11 sys 0m0.001s
12 Test: test02.in
13
14 real 0m0.001s
15 user 0m0.001s
16 sys 0m0.000s
17 Test: test03.in
18
19 real 0m0.001s
20 user 0m0.000s
21 sys 0m0.000s
22 Test: test04.in
23
24 real 0m0.001s
25 user 0m0.001s
26 sys 0m0.000s
27 Test: test05.in
28
29 real 0m0.001s
30 user 0m0.000s
31 sys 0m0.001s
32 Test: test06.in
33
34 real 0m0.001s
35 user 0m0.001s
36 sys 0m0.000s
37 Test: test07.in
38
39 real 0m0.001s
40 user 0m0.001s
41 sys 0m0.000s
42 Test: test08.in
43
44 real 0m0.001s
45 user 0m0.000s
46 sys 0m0.001s
47 Test: test09.in
48
49 real 0m0.001s
50 user 0m0.000s
51 sys 0m0.001s
52 Test: test10.in
53
54 real 0m0.001s
55 user 0m0.001s
56 sys 0m0.000s
57 Test: test11.in
58
59 real 0m0.001s
60 user 0m0.000s
61 sys 0m0.001s
62 Test: test12.in
63
64 real 0m8.163s
65 user 0m8.152s
66 sys 0m0.002s
```

Specifically for `test12.in` it resulted in a run-time greater than 0.4 seconds when compared to the exact same method but using a score array as described above.

```
1 Test      Time
2 sample  0.000010
3 test1   0.000008
4 test2   0.000003
5 test3   0.000004
6 test4   0.000003
7 test5   0.000004
8 test6   0.000004
9 test7   0.000003
10 test8   0.000008
11 test9   0.000010
12 test10  0.000042
13 test11  0.000042
14 test12  7.895575
```

And using `omp_get_wtime()` the total difference is significant between the two implementations (see below).


## Score Array and `pow` Runtime

```
1 ----- TIME TESTS -----
2 Test: sample.in
3 real    0m0.001s
4 user    0m0.000s
5 sys     0m0.001s
6
7 Test: test01.in
8 real    0m0.001s
9 user    0m0.000s
10 sys     0m0.001s
11
12 Test: test02.in
13 real    0m0.001s
14 user    0m0.001s
15 sys     0m0.000s
16
17 Test: test03.in
18 real    0m0.001s
19 user    0m0.000s
20 sys     0m0.001s
21
22 Test: test04.in
23 real    0m0.001s
24 user    0m0.000s
25 sys     0m0.001s
26
27 Test: test05.in
28 real    0m0.001s
29 user    0m0.001s
30 sys     0m0.000s
```

```
31
32 Test: test06.in
33 real    0m0.001s
34 user    0m0.000s
35 sys     0m0.001s
36
37 Test: test07.in
38 real    0m0.001s
39 user    0m0.000s
40 sys     0m0.001s
41
42 Test: test08.in
43 real    0m0.001s
44 user    0m0.001s
45 sys     0m0.000s
46
47 Test: test09.in
48 real    0m0.001s
49 user    0m0.001s
50 sys     0m0.000s
51
52 Test: test10.in
53 real    0m0.001s
54 user    0m0.001s
55 sys     0m0.000s
56
57 Test: test11.in
58 real    0m0.001s
59 user    0m0.000s
60 sys     0m0.001s
61
62 Test: test12.in
63 real    0m7.766s
64 user    0m7.739s
65 sys     0m0.001s
```

When performing a previous project, it came across that a considerable amount of time was being spent in certain library calls, one of them being `pow`. In this case I attempted to discover if there was any impact here. As we can see from the Flame Graph there is a considerable amount of time spent calling the `pow` function and hence it was much more efficient to manually power by multiplying the two values together.





## Without pow Runtime

Due to the issues of `pow` above it was decided to eliminate this and rather multiply as such:

```

1  if (inlist == NULL) {
2      diff += (page_scores[i].score[x] - page_scores[i].score[!x]) *
   (page_scores[i].score[x] - page_scores[i].score[!x]);
3      continue;
4  }
5
6  // Rather than
7      diff += pow(page_scores[i].score[x] - page_scores[i].score[!x], 2);

```

The following times are listed below:

```

1  ----- TIME TESTS -----
2  Test: sample.in
3  real    0m0.001s
4  user    0m0.001s
5  sys     0m0.000s
6
7  Test: test01.in
8  real    0m0.001s
9  user    0m0.001s
10 sys     0m0.000s
11
12 Test: test02.in
13 real    0m0.001s
14 user    0m0.000s
15 sys     0m0.000s
16
17 Test: test03.in
18 real    0m0.001s
19 user    0m0.001s
20 sys     0m0.000s
21
22 Test: test04.in
23 real    0m0.001s
24 user    0m0.000s
25 sys     0m0.000s
26

```




```

27 Test: test05.in
28 real    0m0.001s
29 user    0m0.000s
30 sys     0m0.000s
31
32 Test: test06.in
33 real    0m0.001s
34 user    0m0.000s
35 sys     0m0.001s
36
37 Test: test07.in
38 real    0m0.001s
39 user    0m0.000s
40 sys     0m0.000s
41
42 Test: test08.in
43 real    0m0.001s
44 user    0m0.000s
45 sys     0m0.001s
46
47 Test: test09.in
48 real    0m0.001s
49 user    0m0.001s
50 sys     0m0.000s
51
52 Test: test10.in
53 real    0m0.001s
54 user    0m0.001s
55 sys     0m0.000s
56
57 Test: test11.in
58 real    0m0.001s
59 user    0m0.001s
60 sys     0m0.000s
61
62 Test: test12.in
63 real    0m3.069s
64 user    0m3.064s
65 sys     0m0.001s

```

After modifying the `pow` and using a simple multiplication we see a huge reduction in the time.



The times using `omp_get_wtime()` gave:

1	Dimensions	Time
---	------------	------

2	sample	0.000001
3	test1	0.000001
4	test2	0.000001
5	test3	0.000001
6	test4	0.000001
7	test5	0.000001
8	test6	0.000001
9	test7	0.000001
10	test8	0.000002
11	test9	0.000002
12	test10	0.000014
13	test11	0.000014
14	test12	2.685944

## Loop Unrolling

As an extension in an attempt to optimise the loop overhead, loop unrolling was used on the inner loop:

```

1  for (; i <= npages - 4; i += 4) {
2      page_scores[i].score[x] = dampening_value;
3      diff += update_score(page_scores, &page_scores[i], dampener, x);
4
5      page_scores[i + 1].score[x] = dampening_value;
6      diff += update_score(page_scores, &page_scores[i + 1], dampener, x);
7
8      page_scores[i + 2].score[x] = dampening_value;
9      diff += update_score(page_scores, &page_scores[i + 2], dampener, x);
10
11     page_scores[i + 3].score[x] = dampening_value;
12     diff += update_score(page_scores, &page_scores[i + 3], dampener, x);
13 }
14
15 for (; i < npages; i++) {
16     page_scores[i].score[x] = dampening_value;
17     diff += update_score(page_scores, &page_scores[i], dampener, x);
18 }

```

Although this seemed as if it would provide major performance improvements it only seemed to decrease performance by approximately 0.6 seconds.

1	Test	Time
2	sample	0.000001
3	1	0.000001
4	2	0.000000
5	3	0.000000
6	4	0.000000
7	5	0.000000
8	6	0.000000
9	7	0.000000
10	8	0.000001
11	9	0.000001
12	10	0.000021
13	11	0.000015

This increase in time could possibly be due to the fact that we were making multiple function calls compared to the previous implementation which just performed the whole operation inside the `pagerank` function. Moreover since we are technically increasing the amount of code there is a significant amount of increased usage of the register which contributes to the reduction in performance. Just unrolling the `print` statement after provides no huge benefit as well achieving a time of only 2.727 seconds.

## OpenMP Thread Runtime

```

1  ----- TIME TESTS -----
2  Test: sample.in
3
4  real    0m0.001s
5  user    0m0.000s
6  sys     0m0.001s
7  Test: test01.in
8
9  real    0m0.001s
10 user    0m0.001s
11 sys     0m0.000s
12 Test: test02.in
13
14 real    0m0.001s
15 user    0m0.000s
16 sys     0m0.001s
17 Test: test03.in
18
19 real    0m0.001s
20 user    0m0.001s
21 sys     0m0.000s
22 Test: test04.in
23
24 real    0m0.001s
25 user    0m0.000s
26 sys     0m0.001s
27 Test: test05.in
28
29 real    0m0.001s
30 user    0m0.001s
31 sys     0m0.000s
32 Test: test06.in
33
34 real    0m0.001s
35 user    0m0.000s
36 sys     0m0.001s
37 Test: test07.in
38
39 real    0m0.001s
40 user    0m0.000s
41 sys     0m0.001s
42 Test: test08.in
43

```

```

44 real    0m0.001s
45 user    0m0.000s
46 sys     0m0.003s
47 Test: test09.in
48
49 real    0m0.001s
50 user    0m0.000s
51 sys     0m0.001s
52 Test: test10.in
53
54 real    0m0.001s
55 user    0m0.001s
56 sys     0m0.000s
57 Test: test11.in
58
59 real    0m0.002s
60 user    0m0.003s
61 sys     0m0.003s
62 Test: test12-1.in
63
64 real    0m19.969s
65 user    0m19.910s
66 sys     0m0.014s
67 Test: test12-2.in
68
69 real    0m16.194s
70 user    0m31.585s
71 sys     0m0.051s
72 Test: test12-8.in
73
74 real    0m20.139s
75 user    2m30.563s
76 sys     0m0.111s
77 Test: test12.in
78
79 real    0m14.317s
80 user    0m55.286s
81 sys     0m0.084s

```

The parallel implementation made use of threading the main for loop. What is interesting to note is that a static scheduling seemed like the most efficient approach as the work would easily be divided amongst the threads:

```
#pragma omp parallel for schedule(static)
```

But the idea here is a bit tricky, since the **difference** value is being changed hence the update to difference needs to be extracted.

```

1  #pragma omp parallel for private(i)
2  for (i = 0; i < npages; i++) {
3      page_scores[i].score[x] = dampening_value;
4      register double total = 0.0;
5
6      // Get the list of pages that inlink to this page
7      list* inlist = page_scores[i].page->inlinks;
8
9      // If null then add difference and continue looping

```

```

10     if (inlist == NULL) {
11         continue;
12     }
13     // Get the node to loop
14     node* current = inlist->head;
15     // Loop through the list
16     while (current != NULL) {
17         total += (page_scores[current->page->index].score[!x]) /
18         ((double)current->page-> noutlinks);
19         current = current->next;
20     }
21     page_scores[i].score[x] += total * dampener; // Update the new score
22 }

```

As shown above the difference update is removed into a separate for loop such that there isn't a race condition and eliminates the need for an `atomic` or `critical` section of the for loop as such:

```

1 #pragma omp atomic update
2 diff += ...

```

The following times were achieved for a `default` scheduling:

1	Test	Time
2	sample	0.000013
3	1	0.000013
4	2	0.000001
5	3	0.000001
6	4	0.000001
7	5	0.000018
8	6	0.000016
9	7	0.000016
10	8	0.068941
11	9	0.000003
12	10	0.000026
13	11	0.230591
14	12-1	3.385419
15	12-2	2.316247
16	12-4	2.099469
17	12-8	1.864362

But still these times are nowhere near expected for a parallel optimisation of the pagerank algorithm.

## Pragma Critical Around Difference Sum

Attempting a `pragma omp critical` in a single for loop gives excessively higher runtimes as we can expect multiple threads to be conflicting in contention for this line:

```

1 #pragma omp critical
2 diff = diff + page_scores[i].difference;

```

The times below show the inefficiency of a `pragma omp critical` directive in this scenario:

1	Test	Time
2	sample	0.000015
3	1	0.000016
4	2	0.000001
5	3	0.000001
6	4	0.000001
7	5	0.000017
8	6	0.000015
9	7	0.000018
10	8	0.000032
11	9	0.000002
12	10	0.000024
13	11	0.000153
14	12-1	5.359157
15	12-2	13.226314
16	12-4	23.349664
17	12-8	41.443465

## Pragma Atomic

Since the `pragma omp critical` directive was so time consuming then a more efficient atomic operation was decided upon since the time of contention would be quite small. This means that a **busy-waiting** approach rather than suspending the thread would be much more useful in this scenario:

```
1 #pragma omp atomic write
2 diff = diff + page_scores[i].difference;
```

The results were better than with `critical` yet still not sufficient enough for a parallel solution. But as the number of threads increase to 8 we see an insane projection:

1	Test	Time
2	sample.in	0.000014
3	test01.in	0.000014
4	test02.in	0.000001
5	test03.in	0.000001
6	test04.in	0.000001
7	test05.in	0.000019
8	test06.in	0.000012
9	test07.in	0.000013
10	test08.in	0.066317
11	test09.in	0.000003
12	test10.in	0.000024
13	test11.in	0.461696
14	test12-1.in	4.138385
15	test12-2.in	8.020848
16	test12-4.in	19.682065
17	test12-8.in	76.166786

## Padding Struct

Since some of the issues with time would have been to do with false sharing, the struct was padded to be a total of 64 bytes such that it fits into a cache line directly.

```
1  ----- TIME TESTS -----
2  sample.in    0.000015
3  test01.in    0.000020
4  test02.in    0.000001
5  test03.in    0.000001
6  test04.in    0.000001
7  test05.in    0.000017
8  test06.in    0.000018
9  test07.in    0.000018
10 test08.in    0.000030
11 test09.in    0.000002
12 test10.in    0.000028
13 test11.in    0.000067
14 test12-1.in  3.062184
15 test12-2.in  1.774298
16 test12-8.in  1.606186
17 test12.in    1.704996
18
```

## Reduction on Difference Sum

Instead a reduction on the sum of differences was chosen to help optimise the second for loop:

```
1  #pragma omp parallel for reduction (+:diff)
2  for (i = 0; i < npages; i++) {
3      diff = diff + page_scores[i].difference;
4  }
```

The results with this reduction are very surprising with a consistent decrease every time.

```
1  0.000015
2  0.000015
3  0.000001
4  0.000001
5  0.000001
6  0.000017
7  0.000020
8  0.000015
9  0.000027
10 0.000003
11 0.000025
12 0.000065
13 3.374477
14 2.002390
15 1.850901
```

## Static (Reduction)

With chunk sizes of 2, 4 and 8 we have:

### Chunk Size = 2

```

1  ----- TIME TESTS -----
2  sample.in    0.000004
3  test01.in    0.000004
4  test02.in    0.000001
5  test03.in    0.000001
6  test04.in    0.000000
7  test05.in    0.000001
8  test06.in    0.000004
9  test07.in    0.000004
10 test08.in    0.000007
11 test09.in    0.000002
12 test10.in    0.000021
13 test11.in    0.000048
14 test12-1.in  3.178239
15 test12-2.in  3.051578
16 test12-8.in  2.711370
17 test12.in    2.639630
18

```

### Chunk Size = 4

```

1  ----- TIME TESTS -----
2  sample.in    0.000003
3  test01.in    0.000004
4  test02.in    0.000001
5  test03.in    0.000001
6  test04.in    0.000000
7  test05.in    0.000001
8  test06.in    0.000004
9  test07.in    0.000005
10 test08.in    0.000007
11 test09.in    0.000002
12 test10.in    0.000019
13 test11.in    0.000046
14 test12-1.in  3.309710
15 test12-2.in  3.019476
16 test12-8.in  2.212091
17 test12.in    2.483555

```



### Chunk Size = 8

```
1 0.000033
2 0.000034
3 0.000029
4 0.000027
5 0.000028
6 0.000037
7 0.000026
8 0.000024
9 0.000027
10 0.000023
11 0.000076
12 0.000951
13 0.147611
14 0.000360
15 0.000318
16 0.000345
```

### Dynamic (Reduction)

With the dynamic tests we get a very interesting predicament.

From the data below it seemed as if false sharing was occurring in the loop as we have a sufficiently increased time when using dynamic with a chunk size for 2, 4 and 8 threads. Modifying the chunk sizes to 2, 4, and 8 are shown below:

### Chunk Size = 2

```
1 ----- TIME TESTS -----
2 sample.in 0.000007
3 test01.in 0.000006
4 test02.in 0.000001
5 test03.in 0.000001
6 test04.in 0.000001
7 test05.in 0.000002
8 test06.in 0.000003
9 test07.in 0.000003
10 test08.in 0.157384
11 test09.in 0.000003
12 test10.in 0.000049
13 test11.in 0.000090
14 test12-1.in 6.832125
15 test12-2.in 10.849397
16 test12-8.in 8.730513
17 test12.in 12.916811
```

### Chunk Size = 4

```
1 ----- TIME TESTS -----
2 sample.in 0.000006
3 test01.in 0.000007
4 test02.in 0.000001
```

```

5 | test03.in    0.000001
6 | test04.in    0.000001
7 | test05.in    0.000002
8 | test06.in    0.000003
9 | test07.in    0.000003
10 | test08.in    0.146991
11 | test09.in    0.000003
12 | test10.in    0.000048
13 | test11.in    0.000085
14 | test12-1.in  6.833124
15 | test12-2.in  11.300713
16 | test12-8.in  8.779134
17 | test12.in    12.723585

```

### Chunk Size = 8

However for 8 chunks the results show that there are fewer instances of false sharing:

```

1 | ----- TIME TESTS -----
2 | sample.in    0.000006
3 | test01.in    0.000006
4 | test02.in    0.000001
5 | test03.in    0.000001
6 | test04.in    0.000001
7 | test05.in    0.000002
8 | test06.in    0.000003
9 | test07.in    0.000003
10 | test08.in    0.131622
11 | test09.in    0.000003
12 | test10.in    0.000025
13 | test11.in    0.000061
14 | test12-1.in  3.773181
15 | test12-2.in  3.216287
16 | test12-8.in  2.219597
17 | test12.in    3.090391

```

## Comparison On Number of Pages

As an extension, the methods were run against inputs wherein the number of pages would vary from 100 - 50000000.

